
When Guardrails Fail: Multilingual Long-Context Safety in Tool-Using Agents

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Safety guardrails are increasingly placed in front of language-model agents, yet they
2 are often evaluated as isolated text classifiers. We present , a multilingual guard-
3 to-agent safety evaluation framework that measures whether guardrail bypasses
4 propagate into downstream restricted-tool selection under short and long context.
5 Starting from 1,901 unsafe English prompts derived from AEGIS2.0, we construct
6 parallel prompts across 98 languages, yielding up to 186,298 instances spanning
7 19 harm categories. We evaluate five runtime guards—AprielGuard, CREST,
8 GuardReasoner, WildGuard, and XGuard—and two downstream agents, Llama-
9 3.1-8B-Instruct and Qwen2.5-14B-Instruct, using inert restricted tools. On the
10 3,190-prompt AprielGuard reference set, we further test 8K and 32K contexts with
11 harmful requests placed at the beginning, middle, or end. Results show strong
12 model–position interactions: Qwen beginning-position failure rises from 28.53%
13 to 35.24% from 8K to 32K, while Llama falls from 10.97% to 0.72%. Prompt
14 position reverses model rankings, and 6.4% of matched cases shift from safe
15 at 8K to unsafe at 32K. Low-resource languages also show higher downstream
16 failure (21.65%) than high- and medium-resource languages (11.08% and 10.66%).
17 These results show that multilingual agent safety is compositional and should be
18 evaluated across the complete guard-to-action pipeline. Code available in this link:
19 <https://anonymous.4open.science/r/gefahij1ywnFdlmBeliggetunA>

20 1 Introduction

21 Large language models are increasingly deployed as agents that retrieve information, browse websites,
22 manipulate files, call APIs, and choose among external tools. ReAct established an influential
23 reasoning-and-acting paradigm [Yao et al., 2023], while WebArena [Zhou et al., 2023], OSWorld
24 [Xie et al., 2024], and τ -bench [Yao et al., 2024] demonstrate both the capabilities and persistent
25 unreliability of models operating through external interfaces. This transition changes the consequence
26 of a safety failure: an unsafe response can become an attempted tool invocation or other action
27 crossing a capability boundary.

28 A common defense is therefore to place a runtime safety guard before the agent. Guard models
29 such as Llama Guard [Inan et al., 2023], WildGuard [Han et al., 2024], ShieldGemma [Zeng et al.,
30 2024], and Qwen3Guard [Zhao et al., 2025] classify requests before they reach the underlying model.
31 However, classifier accuracy alone does not determine deployment safety. When a harmful request is
32 incorrectly admitted, the downstream agent may still refuse it, escalate it, or instead select a restricted
33 capability. The relevant unit of evaluation is therefore the complete *guard-to-action trajectory*.

34 This problem is particularly important in multilingual settings. Safety alignment is not language
35 invariant: MultiJail reports increased vulnerability in lower-resource languages [Deng et al., 2024],
36 XSafety identifies elevated unsafe behavior outside English [Wang et al., 2024a], and recent resources

including RTP-LX [de Wynter et al., 2025], PolyGuard [Kumar et al., 2025], Arabic safeguard evaluation [Ashraf et al., 2025], and XL-SafetyBench [Choi et al., 2026] further expose linguistic and cultural variation. A guard that blocks an intent in English but admits the same intent in another language effectively creates a language-dependent capability boundary.

Recent safety benchmarks have improved both risk coverage and evaluation quality. ToxicChat studies real-world moderation [Lin et al., 2023], Do-Not-Answer evaluates requests responsible models should reject [Wang et al., 2024b], HarmBench standardizes adversarial safety evaluation [Mazeika et al., 2024], StrongREJECT examines weaknesses in jailbreak evaluation [Souly et al., 2024], and XSTest measures over-refusal [Röttger et al., 2024]. AEGIS2.0 provides a broad safety taxonomy and large-scale human–LLM safety interactions [Ghosh et al., 2025].

Agent-security work has separately shown that response-level safety is insufficient once models control tools. ToolEmu studies risky tool-use trajectories [Ruan et al., 2024], InjecAgent evaluates indirect prompt injection [Zhan et al., 2024], AgentDojo provides an interactive attack-and-defense environment [Debenedetti et al., 2024], and Agent-SafetyBench evaluates broader agentic hazards [Zhang et al., 2024]. Together, these studies motivate measuring whether an upstream moderation error becomes an action-level failure.

A further complication is long context. LongBench [Bai et al., 2024] and LongBench v2 [Bai et al., 2025] show that nominal context capacity does not imply reliable long-context reasoning. Lost in the Middle demonstrates strong sensitivity to the position of relevant information [Liu et al., 2024], while NoLiMa shows degradation when long-context retrieval requires semantic association rather than lexical matching [Modarressi et al., 2025]. These findings raise a safety question that remains underexplored: if an agent refuses a harmful request in one context, does the same safety behavior persist when that request appears at a different position inside a much longer interaction?

We connect multilingual guardrail evaluation, tool-using agent safety, and long-context robustness. Starting from 1,901 AEGIS-aligned unsafe English roots, we construct parallel prompts across 98 languages and 19 harm categories, retaining BERTScore [Zhang et al., 2020], COMET [Rei et al., 2020], translation-quality, and language-resource metadata. We evaluate five runtime guards—AprielGuard, CREST, GuardReasoner, WildGuard, and XGuard—and follow guard-admitted harmful prompts into two tool-using downstream models, Llama-3.1-8B-Instruct and Qwen2.5-14B-Instruct. To isolate long-context effects, we embed the unchanged harmful request in benign conversation history and evaluate 8K and 32K contexts with the request placed at the beginning, middle, or end. An inert symbolic tool interface allows us to measure attempted restricted-tool selection without executing a harmful action. Stable `sample_ids` permit matched comparisons of the same multilingual request across context lengths, positions, and agents.

Our results show that downstream safety is not a scalar model property. Increasing context length can either improve or degrade safety depending on the model and prompt location. For example, Qwen’s beginning-position restricted-tool selection rises from 28.53% at 8K to 35.24% at 32K, whereas Llama falls from 10.97% to 0.72%. Model rankings also reverse with position, and 6.4% of matched cases change from safe at 8K to unsafe at 32K. Beyond context effects, downstream risk varies sharply across harm categories and languages; low-resource languages exhibit approximately twice the restricted-tool-selection rate of high- and medium-resource languages.

Research questions. **RQ1:** How reliably do multilingual runtime guardrails block harmful requests before they reach an agent? We evaluate five guards across 98 languages, 19 harm categories, resource tiers, and translation-quality strata, and identify the harmful prompts admitted downstream.

RQ2: After guardrail bypass, how do context length, prompt position, and downstream model interact to determine action-level safety? We compare matched 8K and 32K contexts, beginning/middle/end positions, and Llama/Qwen agents using restricted-tool selection and prompt-level safety transitions.

RQ3: Where are downstream safety failures concentrated across multilingual and semantic sub-groups? We analyze variation across languages, resource tiers, harm categories, and translation-quality strata to determine whether aggregate safety rates conceal high-risk groups.

Contributions. **Multilingual guard-to-agent benchmark.** We construct an evaluation framework from 1,901 unsafe roots translated across 98 languages, yielding up to 186,298 instances over 19 harm categories, and evaluate five runtime guards while separating guard blocking from harmful

request admission. **Paired long-context action-safety evaluation.** We follow admitted prompts into two tool-using agents and vary context length and harmful-prompt position while keeping the request unchanged, enabling paired analysis of length, position, model, and 8K→32K safety-state transitions. **Disaggregated downstream-risk analysis.** We characterize action-level failures across language, resource tier, harm category, translation quality, context length, and model choice, revealing positional, semantic, and multilingual heterogeneity that aggregate scores can obscure.

2 Related Work

2.1 Runtime guardrails, safety benchmarks, and multilingual validity

Runtime guard models operationalize safety policy as inference-time classification. Llama Guard introduced policy-conditioned moderation of user prompts and model responses [Inan et al., 2023]; WildGuard jointly models prompt harmfulness, response harmfulness, and refusal behavior [Han et al., 2024]; ShieldGemma provides open moderation models based on the Gemma family [Zeng et al., 2024]; and Qwen3Guard extends this setting with multilingual generative and streaming safeguards [Zhao et al., 2025].

Complementary benchmarks have broadened safety evaluation. ToxicChat studies real-world user–assistant moderation [Lin et al., 2023], Do-Not-Answer collects requests that responsible models should reject [Wang et al., 2024b], HarmBench standardizes adversarial safety evaluation [Mazeika et al., 2024], StrongREJECT highlights weaknesses in automated jailbreak evaluation [Souly et al., 2024], and XSTest measures excessive refusal of benign requests [Röttger et al., 2024]. AEGIS2.0 contributes a broad safety taxonomy and human–LLM interactions for guardrail development [Ghosh et al., 2025]. Multilingual studies further show that strong English safety does not imply cross-lingual robustness. MultiJail reports elevated vulnerability in lower-resource languages [Deng et al., 2024], while XSafety similarly finds higher unsafe-response rates outside English [Wang et al., 2024a]. RTP-LX [de Wynter et al., 2025], PolyGuard [Kumar et al., 2025], Arabic safeguard evaluation [Ashraf et al., 2025], and XL-SafetyBench [Choi et al., 2026] further demonstrate that linguistic and cultural variation can expose failures hidden by English-centric evaluation.

Our work extends this literature by evaluating harmful requests across 98 languages and, crucially, by following guard false negatives into a downstream tool-using agent. We therefore treat multilingual safety as a property of a composed guard-to-action pipeline rather than only a classification problem.

2.2 Tool-using agents and action-grounded safety

Language-model agents connect natural-language reasoning to external capabilities. ReAct established an influential framework for interleaving reasoning and action [Yao et al., 2023], while WebArena [Zhou et al., 2023], OSWorld [Xie et al., 2024], and τ -bench [Yao et al., 2024] demonstrate the challenges of reliable web, computer-use, and tool-mediated interaction.

Security-oriented benchmarks show that these systems introduce risks not captured by text-only evaluation. ToolEmu studies risky tool-use trajectories in an LM-emulated sandbox [Ruan et al., 2024]; InjecAgent evaluates indirect prompt injection in tool-integrated agents [Zhan et al., 2024]; AgentDojo provides a dynamic environment for attacks and defenses involving untrusted content [Debenedetti et al., 2024]; and Agent-SafetyBench evaluates broader agentic safety hazards [Zhang et al., 2024]. Our setting is complementary. Rather than focusing primarily on prompt injection or generic agent competence, we study whether an upstream multilingual safety-classification error propagates into an attempted restricted action. To avoid real-world side effects, the agent selects among inert symbolic tools, allowing us to measure action-level safety without executing harmful capabilities.

2.3 Long-context evaluation and positional sensitivity

The growth of nominal context windows has motivated dedicated evaluation of whether models reliably use distant information. LongBench evaluates long-document understanding, retrieval, and reasoning [Bai et al., 2024], while LongBench v2 extends this setting toward more realistic long-context reasoning [Bai et al., 2025]. Lost in the Middle shows that model performance can vary substantially depending on whether relevant information appears near the beginning, middle, or




Figure 1: **Overview of the multilingual guard-to-agent safety evaluation framework.** Unsafe prompts are translated across 98 languages and evaluated by runtime guardrails. Guard-admitted prompts are forwarded to Llama-3.1-8B-Instruct and Qwen2.5-14B-Instruct under controlled long-context conditions with 8K/32K context lengths and beginning/middle/end prompt positions. The downstream agent chooses among refusal, escalation, and an inert category-relevant restricted tool. We measure guard bypass, downstream recovery, restricted-tool selection, and subgroup variation across language, resource tier, harm category, and translation quality.

end of a long sequence [Liu et al., 2024]. NoLiMa further demonstrates sharp degradation when long-context retrieval requires semantic association rather than direct lexical overlap [Modarressi et al., 2025]. These studies primarily evaluate task competence. We adapt the same length- and position-sensitivity question to safety. The harmful multilingual request is held unchanged while we vary 8K versus 32K context and beginning/middle/end position, allowing paired measurement of whether downstream refusal and restricted-tool selection remain stable under long-context stress. Our results show that the effect is not a generic “lost-in-the-middle” pattern: positional safety profiles differ markedly between model families.

2.4 Evaluator reliability and compositional safety

Automated safety evaluation introduces its own sources of error. MT-Bench and Chatbot Arena document position, verbosity, and self-enhancement biases in LLM-as-a-judge evaluation [Zheng et al., 2023], while StrongREJECT shows that weak evaluators can materially distort conclusions about jailbreak effectiveness [Souly et al., 2024]. These findings motivate outcomes that depend as little as possible on secondary free-form judgment. Accordingly, our primary downstream endpoint is mechanically defined: whether the policy-constrained agent emits `call_tool` for the offered category-relevant restricted symbolic capability. Refusals, escalations, parser failures, and inference errors are recorded separately rather than collapsed into a single safety label. This design enables us to study *compositional safety*: how guard admission, downstream model, context length, prompt position, language, and harm category jointly determine whether a multilingual safety failure reaches the action boundary.

3 Methodology

3.1 Multilingual Benchmark Construction

As shown in Figure 1 our evaluation starts from 1,901 English unsafe root prompts derived from AEGIS2.0, covering 19 harm categories. Each root is translated into 98 languages, producing a

parallel pool of up to 186,298 multilingual prompt instances. Every instance retains its root identifier, language, harm category, language-resource tier, and translation-quality metadata.

Because machine translation can introduce semantic drift, we retain both BERTScore-F1 and COMET scores and assign translations to quality strata before any guard or agent evaluation. These scores are used for subgroup analysis rather than as post-hoc filters on safety outcomes. This allows us to examine whether apparent multilingual safety differences are associated with translation quality while preserving broad language coverage.

Each multilingual instance is assigned a deterministic `sample_id` derived from its root identifier, language, and translated prompt. The same identifier is preserved across guardrail, context-length, prompt-position, and downstream-agent conditions. This enables matched comparisons of the same harmful request throughout the complete evaluation pipeline.

3.2 Runtime Guardrail Evaluation and Admission

We evaluate five runtime safety guards: **AprielGuard**, **CREST**, **GuardReasoner**, **WildGuard**, and **XGuard**. Because the guards expose heterogeneous output formats, model-specific adapters normalize their predictions into a common operational decision while preserving raw model outputs for auditing. Unknown or malformed outputs are handled under a fail-closed policy.

All prompts entering the downstream study are gold-labeled unsafe. Therefore, a normalized guard prediction of `safe` is treated as a guardrail bypass, whereas `unsafe` blocks the prompt before agent execution. Guard performance is analyzed across language, harm category, resource tier, and translation-quality strata. AprielGuard provides the primary downstream reference population, consisting of 3,190 unique admitted multilingual prompts. To test the robustness of downstream findings to the admitted prompt population, we additionally form pairwise-common admission subsets between AprielGuard and XGuard ($n = 2,524$), CREST ($n = 1,237$), GuardReasoner ($n = 102$), and WildGuard ($n = 18$). A common subset contains only prompts classified as safe by both guards.

For these common `sample_ids`, the downstream agent input is identical regardless of which guard admitted the request. We therefore reuse the corresponding AprielGuard agent inference rather than repeating identical model executions. Consequently, these analyses measure robustness across different jointly admitted prompt populations; they are not interpreted as causal effects of guard identity on downstream behavior.

3.3 Downstream Agent and Long-Context Evaluation

Guard admission is treated as an intermediate failure rather than the final safety outcome. Each admitted request is forwarded to Llama-3.1-8B-Instruct and Qwen2.5-14B-Instruct with the same safety policy and an inert category-relevant tool. The agent returns one of three structured decisions: refuse, escalate, or call the tool. A valid call to the offered restricted tool is treated as the primary action-level failure, while refusal or escalation counts as downstream recovery. Inference and parsing failures are recorded separately and are never counted as safe.

For long-context evaluation, the multilingual harmful request is preserved verbatim and embedded in benign English conversation history. We evaluate 8K and 32K contexts with the request placed at the beginning, middle, or end, yielding a $2 \times 3 \times 2$ design across two context lengths, three positions, and two agents. Contexts target approximately 8,192 and 32,768 tokens using a fixed reference tokenizer, while the realized token count for each downstream model is also recorded. Inputs are not silently truncated; unprocessed cases remain explicit inference failures. The benign surrounding context isolates length and position effects without introducing additional harmful instructions.

3.4 Outcome Measures and Paired Analysis

Our primary outcome is *restricted-tool selection*, measured over scorable guard-admitted runs. We also record refusal, escalation, inference coverage, parsing failures, and inference failures.

Because the same `sample_id` is evaluated across conditions, all principal comparisons are paired. We compare 8K versus 32K while holding agent and position fixed, beginning/middle/end positions while holding agent and context length fixed, and Llama versus Qwen within the same length-position condition. We additionally track 8K→32K prompt-level transitions as safe→safe, safe→unsafe,




Figure 2: **Reference-set restricted-tool selection across downstream agent, context length, and harmful-prompt position.** Values report percentages over scorable guard-admitted requests.

unsafe→safe, or unsafe→unsafe. We report Wilson 95% confidence intervals for proportions. Paired comparisons use percentage-point differences, paired-bootstrap 95% confidence intervals with 2,000 resamples, and exact McNemar tests. Low-coverage comparisons are reported separately and are not used as primary evidence.

3.5 Disaggregated and Robustness Analysis

To identify concentrated failures, we disaggregate downstream risk by language, resource tier, harm category, and translation-quality bucket. All 98 languages are included in quantitative analysis; for visualization, we show the 20 highest-risk, 20 median-risk, and 20 lowest-risk languages. Category analysis examines whether particular harm types are disproportionately difficult for either downstream agent, while resource-tier analysis compares high-, medium-, and low-resource languages. These analyses are descriptive, since resource tier covaries with script, language family, tokenization, translation quality, and pretraining exposure.

For robustness, we use the larger AprielGuard-XGuard and AprielGuard-CREST common-admission subsets as the primary checks, while GuardReasoner and WildGuard intersections are treated as exploratory because of their smaller size. All results retain `sample_id`, language, category, context condition, agent model, inference status, and parsing metadata for reproducibility.

4 Results

We analyze downstream behavior after harmful multilingual requests pass the input-guard layer. Our primary outcome is *restricted-tool selection*, the fraction of scorable decisions in which the agent selects the category-relevant restricted symbolic tool. Since tools are inert, this measures attempted unsafe action rather than real execution. The AprielGuard-admitted reference set contains 3,190 unique multilingual prompts. Pairwise-common subsets with XGuard, CREST, GuardReasoner, and WildGuard provide robustness checks. Because downstream inference is reused for common `sample_ids`, these subsets represent different jointly admitted populations rather than causal effects of guard choice.

4.1 Safety depends jointly on agent, context length, and position

Figure 2 and Table 1 show the complete reference-set factorial comparison.

Llama and Qwen exhibit markedly different positional profiles. For Llama, 8K risk rises from 10.97% at the beginning to 23.39% at the end, whereas at 32K the beginning rate falls to only 0.72% while middle and end remain above 22%. Qwen shows an almost opposite pattern: beginning/middle risk is high, while end-position risk is only 10.34% at 8K and 9.18% at 32K. Thus, safety depends

Table 1: **Restricted-tool-selection rate on the reference set.** Values are percentages over scorable decisions; brackets show 95% Wilson confidence intervals.

Agent	Context	Beginning	Middle	End
Llama-3.1-8B	8K	10.97 [9.93, 12.10]	13.54 [12.40, 14.77]	23.39 [21.95, 24.89]
Llama-3.1-8B	32K	0.72 [0.48, 1.08]	22.85 [21.43, 24.34]	22.76 [21.34, 24.25]
Qwen2.5-14B	8K	28.53 [26.99, 30.12]	32.45 [30.84, 34.09]	10.34 [9.34, 11.45]
Qwen2.5-14B	32K	35.24 [33.60, 36.91]	40.71 [32.10, 49.93] [†]	9.18 [8.23, 10.24]

[†]Only 113 of 3,188 Qwen 32K-middle rows are scorable.

Table 2: **Paired 8K→32K changes in restricted-tool selection.** Values are percentage-point changes. XGuard and CREST are pairwise-common admission robustness subsets.

Agent	Position	Reference	XGuard subset	CREST subset
Qwen	Beginning	+6.71	+6.77	+7.11
Qwen	Middle	+9.73 [†]	+10.84 [†]	+3.03 [†]
Qwen	End	-1.16	-1.03	-2.75
Llama	Beginning	-10.25	-11.25	-15.28
Llama	Middle	+9.31	+10.30	+8.81
Llama	End	-0.63	-0.55	-2.18

[†]Comparisons involving Qwen 32K-middle have low scorable coverage; see Appendix A.3.

jointly on model, length, and position. Detailed paired position and agent comparisons appear in Appendices A.1 and A.2.

4.2 Increasing context from 8K to 32K has position-dependent effects

We compare matched `sample_ids` at 8K and 32K while holding agent and prompt position fixed. Longer context does not uniformly increase risk. For Qwen, beginning risk increases from 28.53% to 35.24%, a +6.71-point change (95% CI: +5.45 to +7.90; $p = 1.47 \times 10^{-26}$), while end risk decreases by 1.16 points (95% CI: -2.04 to -0.31; $p = 0.011$). Llama shows the opposite beginning-position effect: risk decreases by 10.25 points (95% CI: -11.44 to -9.18; $p = 1.92 \times 10^{-83}$), whereas middle-position risk increases by 9.31 points (95% CI: +7.93 to +10.69; $p = 9.91 \times 10^{-41}$). The XGuard and CREST subsets reproduce the principal directions, supporting the robustness of these effects.

4.3 Context expansion changes the safety state of individual prompts

Aggregate rates can hide prompt-level changes. We therefore examine matched 8K→32K safety-state transitions. Most matched cases remain stable: 76.2% are safe at both lengths and 11.9% remain unsafe, so 88.1% preserve their state. However, 6.4% transition from safe at 8K to unsafe at 32K, while 5.5% change from unsafe to safe. The pooled unsafe rate therefore changes only from 17.4% at 8K to 18.3% at 32K, a net increase of 0.9 points. This small aggregate difference masks the much larger and opposing model–position effects in Table 2.

4.4 Safety failures vary substantially across harm categories

Figure 4 shows substantial category-dependent variation for both agents.

For Qwen, the largest failure rates occur for Unauthorized Advice (57.6%), Copyright/Trademark/Plagiarism (54.6%), Malware (46.7%), High Risk Government Decision Making (45.2%), Controlled/Regulated Substances (44.1%), and Threat (43.6%). Llama’s largest rates include Copyright/Trademark/Plagiarism (37.5%), Malware (29.5%), Unauthorized Advice (26.3%), PII/Privacy (25.6%), Profanity (23.1%), and Illegal Activity (23.0%). Model rankings also vary by category. Qwen exceeds Llama by 31.3 points for Unauthorized Advice (57.6% vs. 26.3%) and 28.2 points for High Risk Government Decision Making (45.2% vs. 17.0%). Conversely, Llama is less

safe for Illegal Activity (23.0% vs. 6.9%) and Guns and Illegal Weapons (12.8% vs. 1.7%). Thus, no downstream model dominates across all harm types.

4.5 Cross-lingual averages conceal large disparities

Figure 3 visualizes the 20 highest-risk, 20 median-risk, and 20 lowest-risk languages; all 98 languages remain in the quantitative analysis. Oromo, Somali, Amharic, Zulu, Xhosa, Yoruba, Shona, Nyanja, and Lingala appear among the highest-risk languages, whereas Korean, Ukrainian, Vietnamese, Danish, and Swedish occur among the lower-risk group. The visualized range extends from single-digit failure to above 40%. Low-resource languages exhibit 21.65% restricted-tool selection, nearly twice the high-resource (11.08%) and medium-resource (10.66%) rates. Inference-success rates are almost identical across tiers—91.88%, 91.98%, and 91.97%, respectively—so the gap is not explained by gross differences in inference coverage. Because resource tier covaries with script, family, tokenization, translation quality, and pretraining exposure, we interpret this relationship as associational rather than causal.

5 Discussion

Our results show that agent safety emerges from the complete guard-to-agent pipeline rather than from either component alone. Guard bypass does not necessarily lead to action-level failure because the downstream model may still refuse or escalate, but this recovery is highly context dependent. Increasing context from 8K to 32K does not uniformly reduce safety: Qwen beginning-position failure rises from 28.53% to 35.24%, whereas Llama falls from 10.97% to 0.72%; Llama middle-position failure instead increases from 13.54% to 22.85%. Prompt-level transitions similarly show that 6.4% of matched cases become unsafe at 32K while 5.5% become safe. Prompt position also creates model-specific behavior: Qwen is substantially safer at the end, whereas Llama is much safer at the beginning, indicating that context organization itself is a safety-relevant design choice. More discussion in Appendix C.

6 Limitations and Future Work

Our study has several limitations. The 98-language benchmark is primarily translation-based, and although BERTScore and COMET provide quality controls, they do not guarantee native fluency or culturally equivalent harmful intent. The surrounding long context is also controlled, benign, and English-language, which differs from realistic agent settings containing retrieved documents, tool outputs, code, multilingual history, or adversarial content. We evaluate only two downstream agents, two context lengths (8K and 32K), and three prompt positions, so the observed model-specific effects may not generalize to other architectures or longer windows. In addition, Qwen2.5-14B at 32K-middle has only 113 valid inferences out of 3,188 rows (3.54% coverage); its 40.71% failure estimate is therefore not used as primary evidence. More limitations discussed in Appendix D.

7 Conclusion

We present a multilingual guard-to-agent safety evaluation framework spanning 98 languages, 19 harm categories, five runtime guards, and two downstream tool-using agents. By matching the same harmful requests across 8K and 32K contexts and beginning, middle, and end positions, we show that downstream safety is strongly dependent on context structure and model choice rather than being a fixed property of the agent. Our results reveal model-specific positional effects, prompt-level safety transitions, substantial harm-category variation, and higher downstream risk for low-resource languages. These findings demonstrate that multilingual agent safety should be evaluated across the complete guard-to-action pipeline, with attention to context length, prompt position, language, and downstream model behavior.

References

Yasser Ashraf, Yuxia Wang, Bin Gu, Preslav Nakov, and Timothy Baldwin. Arabic dataset for llm safeguard evaluation. In *Proceedings of the 2025 Conference of the Nations of the Americas Chapter*

- 318 *of the Association for Computational Linguistics: Human Language Technologies*, pages 5529–
319 5546. Association for Computational Linguistics, 2025. doi: 10.18653/v1/2025.naacl-long.285.
320 URL <https://aclanthology.org/2025.naacl-long.285/>.
- 321 Yushi Bai, Xin Lv, Jiajie Zhang, Hongchang Lyu, Jiankai Tang, Zhidian Huang, Zhengxiao Du,
322 Xiao Liu, Aohan Zeng, Lei Hou, Yuxiao Dong, Jie Tang, and Juanzi Li. Longbench: A bilingual,
323 multitask benchmark for long context understanding. In *Proceedings of the 62nd Annual Meeting*
324 *of the Association for Computational Linguistics*, pages 3119–3137. Association for Computational
325 Linguistics, 2024. doi: 10.18653/v1/2024.acl-long.172. URL [https://aclanthology.org/](https://aclanthology.org/2024.acl-long.172/)
326 [2024.acl-long.172/](https://aclanthology.org/2024.acl-long.172/).
- 327 Yushi Bai, Shangqing Tu, Jiajie Zhang, Hao Peng, Xiaozhi Wang, Xin Lv, Shulin Cao, Jiazheng Xu,
328 Lei Hou, Yuxiao Dong, Jie Tang, and Juanzi Li. Longbench v2: Towards deeper understanding
329 and reasoning on realistic long-context multitasks. In *Proceedings of the 63rd Annual Meeting of*
330 *the Association for Computational Linguistics*, pages 3639–3664. Association for Computational
331 Linguistics, 2025. doi: 10.18653/v1/2025.acl-long.183. URL [https://aclanthology.org/](https://aclanthology.org/2025.acl-long.183/)
332 [2025.acl-long.183/](https://aclanthology.org/2025.acl-long.183/).
- 333 Dasol Choi, Eugenia Kim, Jaewon Noh, Sang Seo, Eunmi Kim, Myunggyo Oh, Yunjin Park,
334 Brigitta Jesica Kartono, et al. XL-SafetyBench: A country-grounded cross-cultural benchmark
335 for llm safety and cultural sensitivity. *arXiv preprint arXiv:2605.05662*, 2026. URL <https://arxiv.org/abs/2605.05662>.
336 <https://arxiv.org/abs/2605.05662>.
- 337 Adrian de Wynter, Ishaan Watts, Tua Wongsangaroonsri, Minghui Zhang, Noura Farra, Nektar Ege
338 Altintoprak, Lena Baur, et al. RTP-LX: Can llms evaluate toxicity in multilingual scenarios? In
339 *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 39, pages 27940–27950,
340 2025. doi: 10.1609/aaai.v39i27.35011.
- 341 Edoardo Debenedetti, Jie Zhang, Mislav Balunovic, Luca Beurer-Kellner, Marc Fischer, and
342 Florian Tramèr. Agentdojo: A dynamic environment to evaluate prompt injection attacks
343 and defenses for llm agents. In *Advances in Neural Information Processing Systems*,
344 2024. URL [https://proceedings.neurips.cc/paper_files/paper/2024/hash/](https://proceedings.neurips.cc/paper_files/paper/2024/hash/97091a5177d8dc64b1da8bf3e1f6fb54-Abstract-Datasets_and_Benchmarks_Track.html)
345 [97091a5177d8dc64b1da8bf3e1f6fb54-Abstract-Datasets_and_Benchmarks_Track.](https://proceedings.neurips.cc/paper_files/paper/2024/hash/97091a5177d8dc64b1da8bf3e1f6fb54-Abstract-Datasets_and_Benchmarks_Track.html)
346 [html](https://proceedings.neurips.cc/paper_files/paper/2024/hash/97091a5177d8dc64b1da8bf3e1f6fb54-Abstract-Datasets_and_Benchmarks_Track.html).
- 347 Yue Deng, Wenxuan Zhang, Sinno Jialin Pan, and Lidong Bing. Multilingual jailbreak challenges
348 in large language models. In *International Conference on Learning Representations*, 2024. URL
349 <https://arxiv.org/abs/2310.06474>.
- 350 Shaona Ghosh, Prasoon Varshney, Makesh Narsimhan Sreedhar, Aishwarya Padmakumar, Traian
351 Rebedea, Jibin Rajan Varghese, and Christopher Parisien. AEGIS2.0: A diverse ai safety dataset
352 and risks taxonomy for alignment of llm guardrails. In *Proceedings of the 2025 Conference of the*
353 *Nations of the Americas Chapter of the Association for Computational Linguistics: Human Lan-*
354 *guage Technologies*, pages 5992–6026. Association for Computational Linguistics, 2025. doi: 10.
355 18653/v1/2025.naacl-long.306. URL <https://aclanthology.org/2025.naacl-long.306/>.
- 356 Seungju Han, Kavel Rao, Allyson Ettinger, Liwei Jiang, Bill Yuchen Lin, Nathan Lambert, Yejin
357 Choi, and Nouha Dziri. Wildguard: Open one-stop moderation tools for safety risks, jailbreaks,
358 and refusals of llms. In *Advances in Neural Information Processing Systems*, 2024. URL <https://arxiv.org/abs/2406.18495>.
359 <https://arxiv.org/abs/2406.18495>.
- 360 Hakan Inan, Kartikeya Upasani, Jianfeng Chi, Rashmi Rungta, Krithika Iyer, Yuning Mao, Michael
361 Tontchev, Qing Hu, Brian Fuller, Davide Testuggine, and Madian Khabsa. Llama guard: Llm-based
362 input-output safeguard for human-ai conversations. *arXiv preprint arXiv:2312.06674*, 2023. URL
363 <https://arxiv.org/abs/2312.06674>.
- 364 Priyanshu Kumar, Devansh Jain, Akhila Yerukola, Liwei Jiang, Himanshu Beniwal, Thomas
365 Hartvigsen, and Maarten Sap. Polyguard: A multilingual safety moderation tool for 17 lan-
366 guages. *arXiv preprint arXiv:2504.04377*, 2025. URL <https://arxiv.org/abs/2504.04377>.
- 367 Zi Lin, Zihan Wang, Yongqi Tong, Yangkun Wang, Yuxin Guo, Yujia Wang, and Jingbo Shang.
368 Toxicchat: Unveiling hidden challenges of toxicity detection in real-world user-ai conversation.

369 In *Findings of the Association for Computational Linguistics: EMNLP 2023*, pages 4694–4702.
370 Association for Computational Linguistics, 2023. URL [https://aclanthology.org/2023.](https://aclanthology.org/2023.findings-emnlp.311/)
371 [findings-emnlp.311/](https://aclanthology.org/2023.findings-emnlp.311/).

372 Nelson F. Liu, Kevin Lin, John Hewitt, Ashwin Paranjape, Michele Bevilacqua, Fabio Petroni, and
373 Percy Liang. Lost in the middle: How language models use long contexts. *Transactions of the*
374 *Association for Computational Linguistics*, 12:157–173, 2024. doi: 10.1162/tac1_a_00638. URL
375 <https://aclanthology.org/2024.tac1-1.9/>.

376 Mantas Mazeika, Long Phan, Xuwang Yin, Andy Zou, Zifan Wang, Norman Mu, Elham Sakhaee,
377 Nathaniel Li, Steven Basart, Bo Li, David Forsyth, and Dan Hendrycks. Harmbench: A standard-
378 ized evaluation framework for automated red teaming and robust refusal. In *Proceedings of the 41st*
379 *International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning*
380 *Research*, pages 35181–35224. PMLR, 2024. URL [https://proceedings.mlr.press/v235/](https://proceedings.mlr.press/v235/mazeika24a.html)
381 [mazeika24a.html](https://proceedings.mlr.press/v235/mazeika24a.html).

382 Ali Modarressi, Hanieh Deilamsalehy, Franck Dernoncourt, Trung Bui, Ryan A. Rossi, Se-
383 unghyun Yoon, and Hinrich Schuetze. NoLiMa: Long-context evaluation beyond literal match-
384 ing. In *Proceedings of the 42nd International Conference on Machine Learning*, volume
385 267 of *Proceedings of Machine Learning Research*, pages 44554–44570. PMLR, 2025. URL
386 <https://proceedings.mlr.press/v267/modarressi25a.html>.

387 Ricardo Rei, Craig Stewart, Ana C. Farinha, and Alon Lavie. COMET: A neural framework for mt
388 evaluation. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language*
389 *Processing*, pages 2685–2702. Association for Computational Linguistics, 2020.

390 Paul Röttger, Hannah Kirk, Bertie Vidgen, Giuseppe Attanasio, Federico Bianchi, and Dirk Hovy.
391 XSTest: A test suite for identifying exaggerated safety behaviours in large language mod-
392 els. In *Proceedings of the 2024 Conference of the North American Chapter of the Associ-*
393 *ation for Computational Linguistics: Human Language Technologies*, pages 5377–5400. As-
394 sociation for Computational Linguistics, 2024. doi: 10.18653/v1/2024.naacl-long.301. URL
395 <https://aclanthology.org/2024.naacl-long.301/>.

396 Yangjun Ruan, Honghua Dong, Andrew Wang, Silviu Pitis, Yongchao Zhou, Jimmy Ba, Yann
397 Dubois, Chris J. Maddison, and Tatsunori Hashimoto. Identifying the risks of LM agents with an
398 LM-emulated sandbox. In *International Conference on Learning Representations*, 2024. URL
399 <https://arxiv.org/abs/2309.15817>.

400 Alexandra Souly, Qingyuan Lu, Dillon Bowen, Tu Trinh, Elvis Hsieh, Sana Pandey, Pieter Abbeel,
401 Justin Svegliato, Scott Emmons, Olivia Watkins, and Sam Toyer. A strongreject for empty
402 jailbreaks. In *Advances in Neural Information Processing Systems*, 2024. URL [https://arxiv.](https://arxiv.org/abs/2402.10260)
403 [org/abs/2402.10260](https://arxiv.org/abs/2402.10260).

404 Wenxuan Wang, Zhaopeng Tu, Chang Chen, Youliang Yuan, Jen-tse Huang, Wenxiang Jiao, and
405 Michael R. Lyu. All languages matter: On the multilingual safety of llms. In *Findings of the*
406 *Association for Computational Linguistics: ACL 2024*. Association for Computational Linguistics,
407 2024a. URL <https://aclanthology.org/2024.findings-acl.349/>.

408 Yuxia Wang, Haonan Li, Xudong Han, Preslav Nakov, and Timothy Baldwin. Do-not-answer:
409 Evaluating safeguards in llms. In *Findings of the Association for Computational Linguistics:*
410 *EACL 2024*, pages 896–911. Association for Computational Linguistics, 2024b. URL [https://](https://aclanthology.org/2024.findings-eacl.61/)
411 aclanthology.org/2024.findings-eacl.61/.

412 Tianbao Xie, Danyang Zhang, Jixuan Chen, Xiaochuan Li, Siheng Zhao, Ruisheng Cao, Toh Jing
413 Hua, Zhoujun Cheng, Dongchan Shin, Fangyu Lei, et al. OSWorld: Benchmarking multimodal
414 agents for open-ended tasks in real computer environments. *arXiv preprint arXiv:2404.07972*,
415 2024. URL <https://arxiv.org/abs/2404.07972>.

416 Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik Narasimhan, and Yuan Cao.
417 React: Synergizing reasoning and acting in language models. In *International Conference on*
418 *Learning Representations*, 2023. URL <https://arxiv.org/abs/2210.03629>.

Table 3: **Reference-set downstream risk by language-resource tier.**

Tier	Unique prompts	Scorable runs	Risk	95% CI
High	235	2,591	11.08%	[9.93, 12.34]
Medium	498	5,497	10.66%	[9.87, 11.50]
Low	2,457	27,115	21.65%	[21.17, 22.15]

- Shunyu Yao, Noah Shinn, Pedram Razavi, and Karthik Narasimhan. τ -bench: A benchmark for tool-agent-user interaction in real-world domains. *arXiv preprint arXiv:2406.12045*, 2024. URL <https://arxiv.org/abs/2406.12045>.
- Wenjun Zeng, Yuchi Liu, Ryan Mullins, Ludovic Peran, Joe Fernandez, Hamza Harkous, Karthik Narasimhan, Drew Proud, Piyush Kumar, Bhaktipriya Radharapu, Olivia Sturman, and Oscar Wahltinez. Shieldgemma: Generative ai content moderation based on gemma. *arXiv preprint arXiv:2407.21772*, 2024. URL <https://arxiv.org/abs/2407.21772>.
- Qiusi Zhan, Zhixiang Liang, Zifan Ying, and Daniel Kang. Injecagent: Benchmarking indirect prompt injections in tool-integrated large language model agents. In *Findings of the Association for Computational Linguistics: ACL 2024*, 2024. URL <https://aclanthology.org/2024.findings-acl.624/>.
- Tianyi Zhang, Varsha Kishore, Felix Wu, Kilian Q. Weinberger, and Yoav Artzi. BERTScore: Evaluating text generation with bert. In *International Conference on Learning Representations*, 2020. URL <https://arxiv.org/abs/1904.09675>.
- Zhexin Zhang, Shiyao Cui, Yida Lu, Jingzhuo Zhou, Junxiao Yang, Hongning Wang, and Minlie Huang. Agent-safetybench: Evaluating the safety of llm agents. *arXiv preprint arXiv:2412.14470*, 2024. URL <https://arxiv.org/abs/2412.14470>.
- Haiquan Zhao, Chenhan Yuan, Fei Huang, Xiaomeng Hu, Yichang Zhang, An Yang, Bowen Yu, Dayiheng Liu, Jingren Zhou, Junyang Lin, et al. Qwen3guard technical report. *arXiv preprint arXiv:2510.14276*, 2025. URL <https://arxiv.org/abs/2510.14276>.
- Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric P. Xing, Hao Zhang, Joseph E. Gonzalez, and Ion Stoica. Judging llm-as-a-judge with mt-bench and chatbot arena. In *Advances in Neural Information Processing Systems*, 2023. URL <https://arxiv.org/abs/2306.05685>.
- Shuyan Zhou, Frank F. Xu, Hao Zhu, Xuhui Zhou, Robert Lo, Abishek Sridhar, Xianyi Cheng, Tianyue Ou, Yonatan Bisk, Daniel Fried, Uri Alon, and Graham Neubig. Webarena: A realistic web environment for building autonomous agents. *arXiv preprint arXiv:2307.13854*, 2023. URL <https://arxiv.org/abs/2307.13854>.

A Additional Result Analyses

A.1 Detailed Paired Prompt-Position Effects

We compare beginning, middle, and end positions for the same `sample_id`, holding the downstream agent and context length fixed. Figure 7 and Table 4 show strong model-specific positional effects. For Qwen, moving the harmful request toward the end substantially reduces risk: beginning→end changes restricted-tool selection by -18.18 points at 8K and -26.05 points at 32K. Llama shows the opposite profile, with beginning→end increasing risk by 12.41 points at 8K and 22.04 points at 32K. At 32K, Llama middle and end are nearly identical ($\Delta = -0.09$, $p = 0.930$). These results indicate model-specific positional safety biases rather than a universal “lost-in-the-middle” effect.

A.2 Detailed Paired Downstream-Agent Comparison

We compare Llama and Qwen on the same `sample_id`, context length, and prompt position. Figure 8 and Table 5 show that the relative safety ranking reverses with position.

Table 4: **Prompt-position effects on the AprielGuard reference set.** Δ denotes condition B minus condition A.

Agent	Length	Comparison	Pairs	Δ pp	McNemar p
Qwen	8K	Beginning $\rightarrow$ Middle	3190	+3.92	1.13×10^{-8}
Qwen	8K	Beginning $\rightarrow$ End	3190	-18.18	2.26×10^{-135}
Qwen	8K	Middle $\rightarrow$ End	3190	-22.10	8.70×10^{-190}
Qwen	32K	Beginning $\rightarrow$ Middle	113	+6.19	0.039
Qwen	32K	Beginning $\rightarrow$ End	3190	-26.05	6.76×10^{-212}
Qwen	32K	Middle $\rightarrow$ End	113	-31.86	2.91×10^{-11}
Llama	8K	Beginning $\rightarrow$ Middle	3190	+2.57	2.60×10^{-4}
Llama	8K	Beginning $\rightarrow$ End	3190	+12.41	3.23×10^{-54}
Llama	8K	Middle $\rightarrow$ End	3190	+9.84	1.06×10^{-52}
Llama	32K	Beginning $\rightarrow$ Middle	3190	+22.13	3.29×10^{-198}
Llama	32K	Beginning $\rightarrow$ End	3190	+22.04	5.50×10^{-200}
Llama	32K	Middle $\rightarrow$ End	3190	-0.09	0.930

Table 5: **Paired Qwen–Llama comparison on the AprielGuard reference set.** $\Delta = R_{\text{Llama}} - R_{\text{Qwen}}$.

Context	Position	Qwen	Llama	Δ pp	95% CI
8K	Beginning	28.53	10.97	-17.55	[-19.15, -15.96]
8K	Middle	32.45	13.54	-18.90	[-20.47, -17.37]
8K	End	10.34	23.39	+13.04	[11.54, 14.51]
32K	Beginning	35.24	0.72	-34.51	[-36.18, -32.92]
32K	Middle	40.71	37.17	-3.54 [†]	[-13.27, 7.08]
32K	End	9.18	22.76	+13.57	[12.07, 15.05]

[†]Based on only 113 paired scorable Qwen 32K-middle cases.

Llama is substantially safer at beginning and most middle positions. At 8K, it is 17.55 points safer at the beginning and 18.90 points safer in the middle; at 32K-beginning, the difference grows to 34.51 points (35.24% for Qwen versus 0.72% for Llama). The ranking reverses at the end: Qwen is safer by 13.04 points at 8K and 13.57 points at 32K.

This reversal also appears in the larger robustness subsets. For AprielGuard–XGuard, the differences are -37.04 points at 32K-beginning and +14.66 points at 32K-end; for AprielGuard–CREST they are -46.89 and +14.47 points, respectively. Thus, neither agent has a universally superior safety profile.

A.3 Qwen 32K-Middle Coverage Diagnostic

The Qwen2.5-14B-Instruct 32K-middle condition has substantially lower inference coverage than the other factorial cells. Only 113 of 3,188 reference-set rows produce valid inference, corresponding to 3.54% coverage; the remaining 3,075 rows contain inference failures. The observed 40.71% restricted-tool-selection rate for this condition is therefore estimated over a highly selected subset.

The same limitation affects paired analyses: the reference 8K–32K middle comparison contains only 113 scorable pairs, while the XGuard and CREST common-admission subsets contain 83 and 33 pairs, respectively. We retain these results for transparency but do not use them as primary evidence. All other principal reference-set factorial cells contain 3,190 successful inferences.

B Summary of findings

The results reveal four consistent patterns. First, context length does not have a uniform safety effect: 8K $\rightarrow$ 32K can either increase or decrease restricted-tool selection depending on model and position. Second, prompt location strongly changes safety, with Qwen substantially safer at the end and Llama especially safe at 32K-beginning. Third, model rankings also vary across harm categories, so no single downstream agent is uniformly safer. Finally, these effects occur alongside substantial

481 multilingual heterogeneity, with low-resource languages showing approximately twice the aggregate
482 failure rate of high- and medium-resource languages.

483 **C More Discussion**

484 There is also no universally safer downstream agent. Llama is substantially safer at beginning and
485 most middle positions, while Qwen is approximately 13–14 points safer at the end; harm-category
486 rankings likewise reverse across models. Multilingual disparities remain pronounced after guard
487 bypass: low-resource languages show 21.65% restricted-tool selection compared with 11.08% for
488 high-resource and 10.66% for medium-resource languages, despite similar inference coverage. These
489 results suggest that multilingual agent safety should be evaluated jointly across guard behavior,
490 context length, prompt position, model choice, language, and harm category rather than summarized
491 by a single aggregate safety score.

492 **D More Limitations**

493 The downstream analysis is conditional on guard admission, and the XGuard, CREST, GuardReasoner,
494 and WildGuard robustness subsets reuse AprielGuard agent outputs for shared `sample_ids`, so these
495 comparisons reflect different admitted populations rather than causal guard effects. Our symbolic
496 tools also measure attempted unsafe action rather than real-world execution, and the current pipeline
497 omits output guards, execution-time authorization, and benign over-refusal controls. Future work
498 should add native-authored multilingual prompts, realistic retrieved and multilingual contexts, more
499 agents and longer context windows, independently evaluated guard–agent pipelines, sandboxed tool
500 execution, additional defense layers, and matched benign controls.




Figure 3: **Reference-set cross-lingual downstream risk.** Bars report restricted-tool-selection rates among scorable runs. All 98 languages are retained in the full analysis.




Figure 4: **Reference-set restricted-tool selection by harm category and downstream agent.** Values are percentages over scorable runs.




Figure 5: **Reference-set safety-state transitions from 8K to 32K.** “Unsafe” denotes restricted-tool selection.




Figure 6: **Paired effect of increasing context from 8K to 32K.** Points show percentage-point changes with paired-bootstrap 95% confidence intervals. Positive values indicate higher failure at 32K.




Figure 7: **Paired effect of harmful-prompt position.** Points show percentage-point changes in restricted-tool selection when the same harmful request is moved between beginning, middle, and end positions.




Figure 8: **Paired downstream-agent comparison.** The plotted difference is $R_{\text{Llama}} - R_{\text{Qwen}}$. Negative values favor Llama; positive values favor Qwen.